





```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.impute import SimpleImputer
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Flatten, Dense
from sklearn.metrics import accuracy_score, confusion_matrix
import joblib

# Load the datasets
train_data = pd.read_csv(r'mnist_train.csv')
test_data = pd.read_csv(r'mnist_test.csv')

# Explore the dataset
print("Shape of Training Data:", train_data.shape)
print("\nShape of Test Data:", test_data.shape)

# Identify unique classes
unique_classes = train_data['label'].unique()
num_classes = len(unique_classes)
print("\nNumber of Unique Classes:", num_classes)
```



```
# Extract features and labels
X_train = train_data.iloc[:, 1:] / 255.0 # Normalize pixel values
y_train = train_data.iloc[:, 0]

X_test = test_data.iloc[:, 1:] / 255.0
y_test = test_data.iloc[:, 0]

num_features = X_train.shape[1]
print("\nNumber of features: ", num_features)

print("-----")

# Check for missing values in the training dataset
missing_values_train = X_train.isnull().sum()
print("\nMissing values in X_train:\n", missing_values_train)

print("-----")

# Check for missing values in the test dataset
missing_values_test = X_test.isnull().sum()
print("\nMissing values in X_test:\n", missing_values_test)

print("-----")

# Reshape images
image_shape = (28, 28)
X_train = X_train.values.reshape(-1, *image_shape) # Reshape to 28x28
X_test = X_test.values.reshape(-1, *image_shape)
```

```
# Visualize some images from the training dataset
plt.figure(figsize=(10, 10))
for i in range(9):
    plt.subplot(3, 3, i + 1)
    plt.imshow(X_train[i], cmap='gray')
    plt.title(f"Label: {y_train.iloc[i]}")
    plt.axis('off')
print('\nVisualize some images from the dataset \n')
plt.show()

# Split dataset into train and validation sets
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.2, random_state=42)

print("-----")

# Print shapes of processed datasets
print("\nShapes after Preprocessing:")
print("X_train:", X_train.shape)
print("y_train:", y_train.shape)
print("X_val:", X_val.shape)
print("y_val:", y_val.shape)
print("X_test:", X_test.shape)
print("y_test:", y_test.shape)
print("-----")
```

```
# Handle missing values using an imputer transformer
imputer = SimpleImputer(strategy='mean')
X_train = imputer.fit_transform(X_train.reshape(X_train.shape[0], -1))
X_val = imputer.transform(X_val.reshape(X_val.shape[0], -1))
X_test = imputer.transform(X_test.reshape(X_test.shape[0], -1))

# Initial Experiment: K-NN classification
knn = KNeighborsClassifier()
param_grid = {'n_neighbors': [3, 5, 7]}
grid_knn = GridSearchCV(knn, param_grid, cv=5)
grid_knn.fit(X_train, y_train)
best_knn = grid_knn.best_estimator_
knn_predictions = best_knn.predict(X_val)
knn_accuracy = accuracy_score(y_val, knn_predictions)
print("\nK-NN Accuracy on Validation Data:", knn_accuracy)

print("-----")

# Subsequent Experiment: ANN
num_hidden_neurons = [64, 128]
learning_rates = [0.001, 0.01]
batch_sizes = [32, 64]
best_ann_accuracy = 0.0
best_ann_model = None
```

```

for hidden_neurons in num_hidden_neurons:
    for learning_rate in learning_rates:
        for batch_size in batch_sizes:
            ann_model = Sequential()
            ann_model.add(Dense(hidden_neurons, activation='relu', input_shape=(784,)))
            ann_model.add(Dense(num_classes, activation='softmax'))

            ann_model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=learning_rate),
                              loss='sparse_categorical_crossentropy', metrics=['accuracy'])

            ann_model.fit(X_train, y_train, batch_size=batch_size, epochs=10, verbose=0)
            ann_predictions = np.argmax(ann_model.predict(X_val), axis=-1)
            ann_accuracy = accuracy_score(y_val, ann_predictions)

            print(f"ANN Accuracy (hidden neurons={hidden_neurons}, learning rate={learning_rate}, "
                  f"batch size={batch_size}): {ann_accuracy}")

            if ann_accuracy > best_ann_accuracy:
                best_ann_accuracy = ann_accuracy
                best_ann_model = ann_model

print("-----")

# Compare outcomes of experiments
print('\nComparison of Outcomes:')
print("K-NN Accuracy on Validation Data:", knn_accuracy)
print("Best ANN Accuracy on Validation Data:", best_ann_accuracy)

```

Activate  
Go to Settings

```

# Get confusion matrix of the best model
best_ann_predictions = np.argmax(best_ann_model.predict(X_val), axis=-1)
confusion_mat = confusion_matrix(y_val, best_ann_predictions)
print("\nConfusion Matrix:")
print(confusion_mat)
print("-----")

# Save the best model
joblib.dump(best_ann_model, 'best_ann_model.pkl')

# Reload the best model from a separate file
loaded_model = joblib.load('best_ann_model.pkl')

# Apply the best model on the testing data
test_predictions = np.argmax(loaded_model.predict(X_test), axis=-1)
test_accuracy = accuracy_score(y_test, test_predictions)
print("\nTesting Accuracy using the Best Model:", test_accuracy)

```



Shape of Training Data: (55086, 785)



Shape of Test Data: (10000, 785)

Number of Unique Classes: 10

Number of features: 784

-----  
Missing values in X\_train:

1x1 0

1x2 0

1x3 0

1x4 0

1x5 0

..

28x24 1

28x25 1

28x26 1

28x27 1

28x28 1

Length: 784, dtype: int64

-----  
Missing values in X\_test:

1x1 0

1x2 0

1x3 0

1x4 0

1x5 0

..

28x24 0

28x25 0

28x26 0

28x27 0


28x28 0

Length: 784, dtype: int64


-----

Visualize some images from the dataset


Label: 5




Label: 0




Label: 4




Label: 1



Label: 9




Label: 2




Activate Windows  
Go to Settings to activate Windows.

Label: 1




Label: 9



Label: 2




Label: 1



Label: 3



Label: 1



Activate Windows  
Go to Settings to activate Windows.



Shapes after Preprocessing:

X\_train: (44068, 28, 28)

y\_train: (44068,)

X\_val: (11018, 28, 28)

y\_val: (11018,)

X\_test: (10000, 28, 28)

y\_test: (10000,)

-----  
K-NN Accuracy on Validation Data: 0.9694136866944999

-----  
345/345 [=====] - 1s 2ms/step

ANN Accuracy (hidden neurons=64, learning rate=0.001, batch size=32): 0.9673261935015429

345/345 [=====] - 1s 2ms/step

ANN Accuracy (hidden neurons=64, learning rate=0.001, batch size=64): 0.963605009983663

345/345 [=====] - 1s 2ms/step

ANN Accuracy (hidden neurons=64, learning rate=0.01, batch size=32): 0.9513523325467417

345/345 [=====] - 1s 2ms/step

ANN Accuracy (hidden neurons=64, learning rate=0.01, batch size=64): 0.9609729533490652

345/345 [=====] - 1s 2ms/step

ANN Accuracy (hidden neurons=128, learning rate=0.001, batch size=32): 0.9721365039027047

345/345 [=====] - 1s 2ms/step

ANN Accuracy (hidden neurons=128, learning rate=0.001, batch size=64): 0.9732256307859866

345/345 [=====] - 1s 2ms/step

ANN Accuracy (hidden neurons=128, learning rate=0.01, batch size=32): 0.9576148121256126

345/345 [=====] - 1s 2ms/step

ANN Accuracy (hidden neurons=128, learning rate=0.01, batch size=64): 0.9601561081866037

-----

Comparison of Outcomes:

K-NN Accuracy on Validation Data: 0.9694136866944999

Best ANN Accuracy on Validation Data: 0.9732256307859866

-----  
345/345 [=====] - 1s 2ms/step

Confusion Matrix:

```
[[1121  0  2  1  3  1  0  1  0  3]
 [ 1 1173  2  1  2  0  2  3  0  0]
 [ 2  1 1059  3  0  1  4 11  1  2]
 [ 1  1 15 1098  1  9  0 10  7  5]
 [ 1  0  1  1 1031  0  2  3  0 13]
 [ 4  2  0 12  1 981  4  3  1  0]
 [ 4  1  1  0  2 10 1086  0  2  0]
 [ 0  5 10  0  4  0  0 1151  0  5]
 [ 4  8  7 13  4 10  5  5 965  9]
 [ 5  2  0  3 16  4  0 11  1 1058]]
```

-----  
313/313 [=====] - 1s 2ms/step

Testing Accuracy using the Best Model: 0.976

